N83 K = 16,627: root-propagation discriminator

Run: 2026-10-18 UTC. **Curve:** `icv1-f2m83-tm6151469093347-debefd74`. The frozen protocol is `PROTOCOL.md`. The exact retained signed-Frobenius base has 16,627 columns and 2,760,082 points. The accepted native worker source is clean pre-stack commit `43c9850a73f1d960f65ed470fdb8adac981a5ffe`; `source-files.sha256`, `inputs-before-run.sha256`, `BUILD_RECEIPT.md`, `VALIDATION.md` and `COMMANDS.md` retain source, object, binary, build and test checks. The complete source-linked account is `RUN_REPORT.md`.

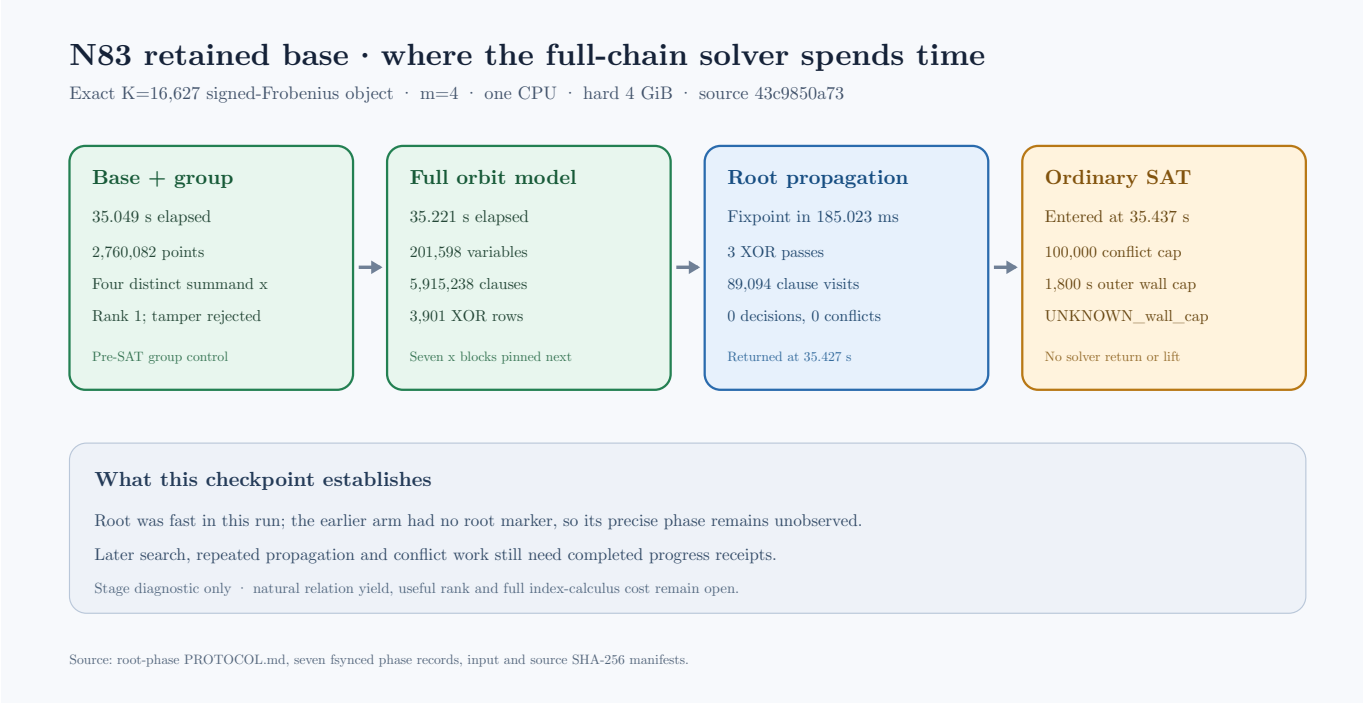

Frozen positive arm

Measure	Observed value
Status	`UNKNOWN_wall_cap`; exit 124; no OOM
UTC interval	05:22:09–05:52:09
CPU / hard memory / swap	1 / 4 GiB / 0
Cgroup peak	2,661,863,424 bytes
Retained object	83,701,045 compressed bytes; 16,627 orbit columns
Pre-SAT group check	rank 1; tampered target rejected
Full m=4 model	201,598 variables; 5,915,238 clauses; 3,901 XOR rows
Root propagation	`Fixpoint` in 185.023 ms; zero decisions and conflicts
Ordinary SAT entry	35.437 s elapsed; no solver return by outer cap

Interpretation

The initial level-zero root pass returned on this exact seven-block witness in the instrumented run. That makes it an unlikely explanation for the preceding arm's cap, but the earlier arm had no root marker. A later decision, propagation, conflict or restart path still needs completed progress receipts. Post-root decisions, conflicts and models remain unknown. The pre-SAT rank-one check is a planted group preflight, not a returned SAT relation or natural-query useful rank. A cap is an unresolved bounded outcome; this run supplies no factor-base winner, complete index-calculus cost or matched-rho ratio.

Raw source: `positive/worker.json.phases.jsonl`, `positive/container-state.json` and `positive/outcome.json`. Stdout/stderr, cgroup peak and the exact UTC interval are retained beside them. Failed worker builds remain linked from `BUILD_RECEIPT.md`.