Cheap degree-two refutations: local work screen

Source 9bed104b0a. Same registered curves, standard bases and public target0/round1 as the full panel; one warmup and one measured round, four arms per curve. All16 executions verify, including8 warmups; all8 measured replays reproduce. Gate passed: True. The native replay-all mode excludes warmups. This corrects earlier prose equating total verified records with measured replays. No L0 elapsed-time or pipeline speedup claim follows.

The failed all-width degree-three policy spends3.000 billion word XORs on the n23 F4 target; the guarded policy spends62.421 million, against79.418 million baseline. The old native full-pipeline regression remains in its own report; logical work counts here are not native wall time.

Both original16/22-variable fixtures retain identical complete roots and all checked deterministic counters against source71c4c99d85. The final corrected policy receives its own identity; pre-version screen records keep their sealed bytes and frozen-source replay. Sources: the adjacent Rust-prepared summary, ADAPTIVE_PROTOCOL.md and the raw screen tree.

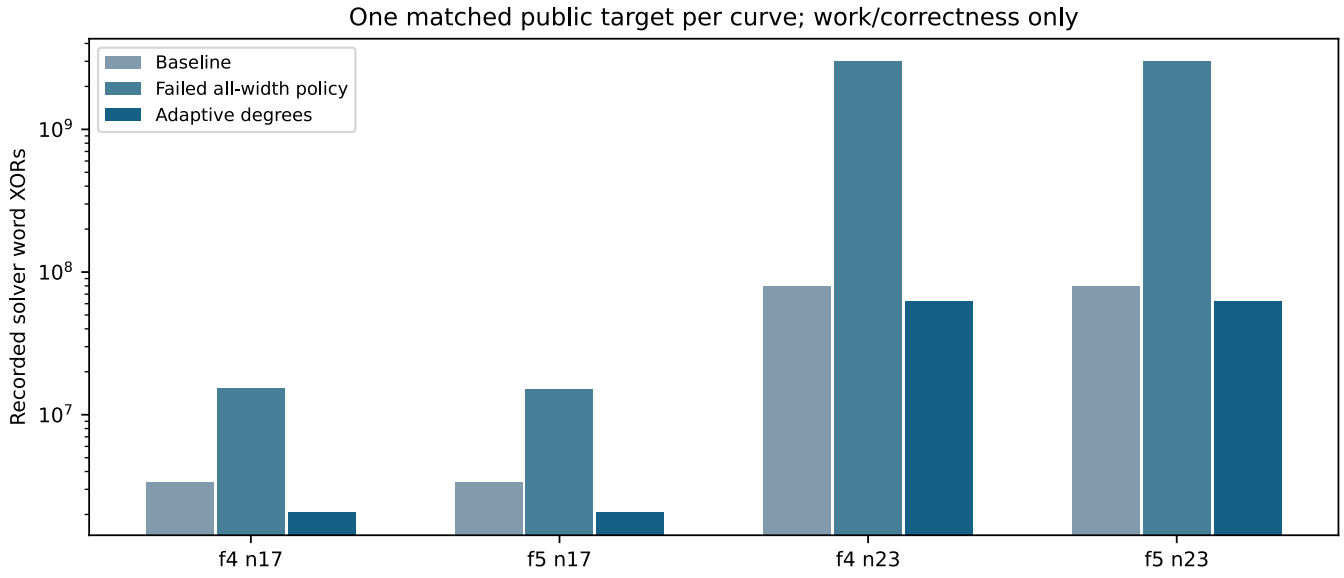

Matched target work

n	Engine	Calls	Baseline XORs	Failed XORs	Adaptive XORs	Queries/answer equal
17	f4	2398	3355654	15190651	2062875	True
17	f5	2398	3350611	15107801	2058669	True
23	f4	135075	79417986	3000248688	62421379	True
23	f5	135075	79343067	2994993068	62338849	True