Four follow-ups: cross-lead review packet

Frozen public 21-model, 1,260-target cohort; 22 October 2026

Accounting. Full cold candidate cost divided by same-process matched negation-rho cost at $M + 0.5S + 10I$ and 0.1 source-addition equivalents per native logical operation. The external column uses the archived wrapped graph collector. Values are ratios of summed costs across the stated targets; the accompanying SVG shows per-size ratios. All 36 price cells and exact source grid paths are in `results/standard-cost-grid.csv`.

Arm	Candidate	Rho	Candidate/rho	Candidate/external
Full base, 16 jumps	1,260	1,260	2.448830	1.000000
Full base, 8 jumps	1,260	1,260	2.447173	0.999323
Full base, 32 jumps	1,260	1,260	2.453278	1.001816
Half base, 16 jumps	1,260	1,260	3.017732	1.232316
Quarter base, 16 jumps	1,260	1,260	3.919193	1.600435
Projective wrapper	1,260	1,260	3.793783	1.549223
Retained weighted state	1,260	1,260	4.179297	1.706651

Validation. Factor-base arms have 6,300 candidate and 6,300 rho source-scalar checks. Each projective-law arm has 1,260 candidate and 1,260 rho source-scalar checks. Its independent Rust validator recomputed 576 price cells. The weighted-law arithmetic gate includes targeted exceptional cases, but finite checks do not prove universal formula completeness. Study-specific raw JSONL, validators, manifests, and hashes remain in their source directories. **Point-driven pencil.** The validated sparse-rank continuation has 543 distinct candidate solves and 544 source-verified rho solves. One 32-bit target reached caps under both online and peel; both candidate costs and ratios are null. Another 717 original targets have no completed candidate solve. Nine 60-target curves and one 28-bit single target have all 36 paired online/peel/rho prices. The next table retains the representative prices and the number of cells where peel costs less than online; the companion pencil SVG shows all-cell ranges.

p	Targets	Peel/online	Peel/rho	Cells peel<online
739	60	0.898244	11.615357	24/36
1,151	60	0.796603	22.836038	24/36
1,187	60	0.736053	22.877204	24/36
3,371	60	0.379821	93.480568	24/36
3,527	60	0.393296	105.157423	24/36
3,931	60	0.365357	121.367036	24/36
11,731	60	0.278124	1,204.654657	24/36
11,887	60	0.282829	1,090.655477	24/36
12,143	60	0.269806	1,146.700226	24/36
48,383	1	0.233555	22,644.843832	24/36

Every paired curve also has 12/36 cells where peel costs more than online and 0/36 where peel costs less than rho. The 28-bit row is one target, not a cohort mean. Its 36-cell range and source paths are in the price grid; all 1,448 pencil arm-target rows, including two unpriced failures, are in the target index. **Kummer open requirement.** The coordinate gate checked 8,820 labeled cases on all 1,260 public targets: 8,813 matching nonzero xADD outputs, seven recorded zero outputs, and no wrong nonzero outputs. Its separate signed-Recover gate repaired 175 quotient zero cases and recovered all 1,260 signed targets (1,257 generic, three bounded exceptional). The priced labeled DLP collector remains open, so its candidate cost, rho cost, and ratio are null here. **Review files.** `REPORT.md` links each source study and its raw data. `DATA_DICTIONARY.md` defines both normalized CSV files, empty fields, and exact source references. The target index has 11,528 rows sharing 1,260 checked workload keys and retains all seven Kummer label outcomes per target. The two SVGs show factor-base/law and paired pencil prices. The CSVs and summary JSON support independent regrouping without using this report's wording.