CPU recovery ladder: scheduling deadline

Observations only. This attempt produced no larger-field recovery evidence. The existing completed CPU comparison is separate and remains valid within its own fixed fixtures and source identity.

Job `hashed-ecdl-cpu-scale-20261030` was submitted October 29, 2026 at 20:34:24 PDT and reached its 1500-second deadline at 20:59:24 PDT. The last observed Pod state was Pending/Unschedulable: the available CPU host was occupied by another experiment with exclusive-host scheduling. The other workload was left intact. The terminal job condition is `DeadlineExceeded`; no recovery rows, certificates, workload logs or result archive were produced. No automatic retry was submitted.

Bounded CPU scaling attempt: no recovery observations

No solver rows, certificates, or algorithm durations; no quantitative scaling plot.

Job: hashed-ecdl-cpu-scale-20261030; terminal receipt: scale/results/status.json

The frozen plan contains 24 executions: six exact prime-field curves from 20 through 40 bits, two public point searches per curve, and two identical CPU arm invocations. Each completed recovery would have been replayed in a separate process and would have rejected a corrupted scalar. The intended limits were 120 seconds per recovery,60 seconds per external replay,20 million retained tuple rows and 16384 columns. One job reserved 2 vCPUs and 4 GiB. These are requested bounds, not measured memory usage or performance.

Targets are constructed by searching public x coordinates from fixed seeds 101 and 202, without constructing a target logarithm. The source uses eight ordered factor-base slots with tau equal to twice the integer ceiling of r^(1/8), plus 2. The native executable and independent archive verifier were compiled and hashed before submission. The standalone source is `a09a3b1bf0af`; the runtime image is pinned separately to the main-built image recorded in `scale/results/build.json`.

Protocol SHA-256: `91e3449df0ae301d9d882db910f334a1e4ffa9e93cbd1ba93630f9d3b16f7510`. Job UID: `2c2bb11e-f38b-47b5-9e9e-3e069fbbfa24`. The terminal receipt and incomplete export are `scale/results/status.json` and `scale/results/export.json`. Protocol plan: 24 cells; export contains 0 execution rows and 0 completions. The strict completeness checker rejects the empty export against that plan; `scale/results/export-check.log` preserves its findings. No rows were invented. The native analyzer is identified in the frozen contract but did not run.

The 1500 seconds are scheduling elapsed time, not algorithm runtime. No field-size frontier or largest feasible CPU scale can be inferred. A fresh bounded run needs available CPU capacity and its own immutable run record. The required larger-scale validation remains open; the implementation and fixed-workload comparison are reviewable separately.