Matched accounting for the factored source

2026-11-03 · T-1576/T-1579

Baseline/A-A use the previous homogeneous consumer; candidate factors seed dependence into two final powers. Reversible instrumentation covers residual products, final powers and four separate public-exponent setup primitives. Six non-source modules are byte-identical.

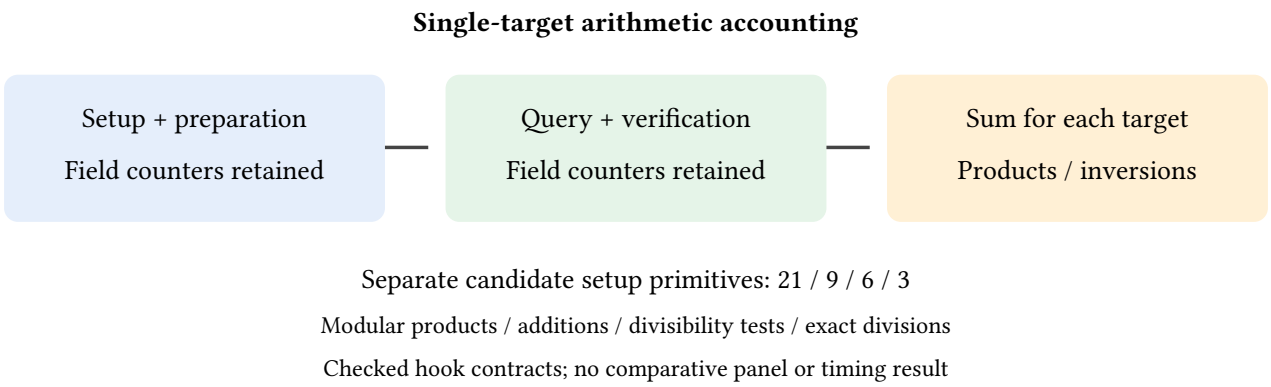

Checked instrument

Both optimized binaries compile. Two selected native tests pass,84 inherited tests filtered, covering all8 complete recovery receipts and setup/stage contracts. The candidate exponent calculation makes21 modular-product calls,9 modular-add calls,6 divisibility tests and3 exact integer divisions. These categories are separate from field arithmetic and do not price every machine operation.14 source maps invert exactly;72 prebuild source/data locks verify.

Analysis

The new21-slot schemas prevent confusion with historical17-slot rows. Seven synthetic strict-parser tests pass. Two additional tests verify single-target aggregation: added setup can reverse a query win, setup inversions affect the symbolic break-even weight, and A/A failures remain invalid. The primary field-count verdict uses setup plus preparation plus query plus verification for each target. Query-only results and integer setup counters remain separate. This arithmetic sum is not cold elapsed time.

Scope and next action

No comparative panel or cloud job ran. The separately bound draft remains unfrozen pending new-instrument preflight, quota and charged/reserved allowance evidence. Prior panels and original4563-case cohort remain intact. Next execute the admitted comparison and retain all stage costs. Evidence: source-lock.json, instrumentation maps, receipt.json, raw logs, analysis-lock.json and comparison-protocol.draft.json. Prior context: the checked seed-factor derivation, native accounting definitions and canonical experiment-protocol instructions.