Check the target coordinate, not just full-system rank

The post-v8 diagnostic verifies all18 unchanged relation systems. The ordinary base fixes the target log on all nine models; the twist-progression base leaves it ambiguous on all nine. Native RREF and independent PARI `matker` and `matinverseimage` agree. No new decomposition query or changed base is used.

For each mod823 matrix, the native program retains a particular solution and full nullspace basis and verifies every affine row/direction. A target is identifiable exactly when every nullspace vector has zero target entry. Thus rank deficiency alone is not used as evidence of failed target recovery. Separate controls cover deficient matrices with identifiable and ambiguous target coordinates.

The ordinary system has rank9 for9 unknowns and verifies k37 on public points. The twist system has rank8, free column8 (the target), particular solution `[1,39,0,29,2,809,739,21,0]` and nullspace direction `[0,3,822,3,822,820,816,1,1]`. Adding a multiple of that direction changes k while satisfying every recorded relation. Those rows alone therefore do not determine k. Both patterns persist exactly under every certified isogeny.

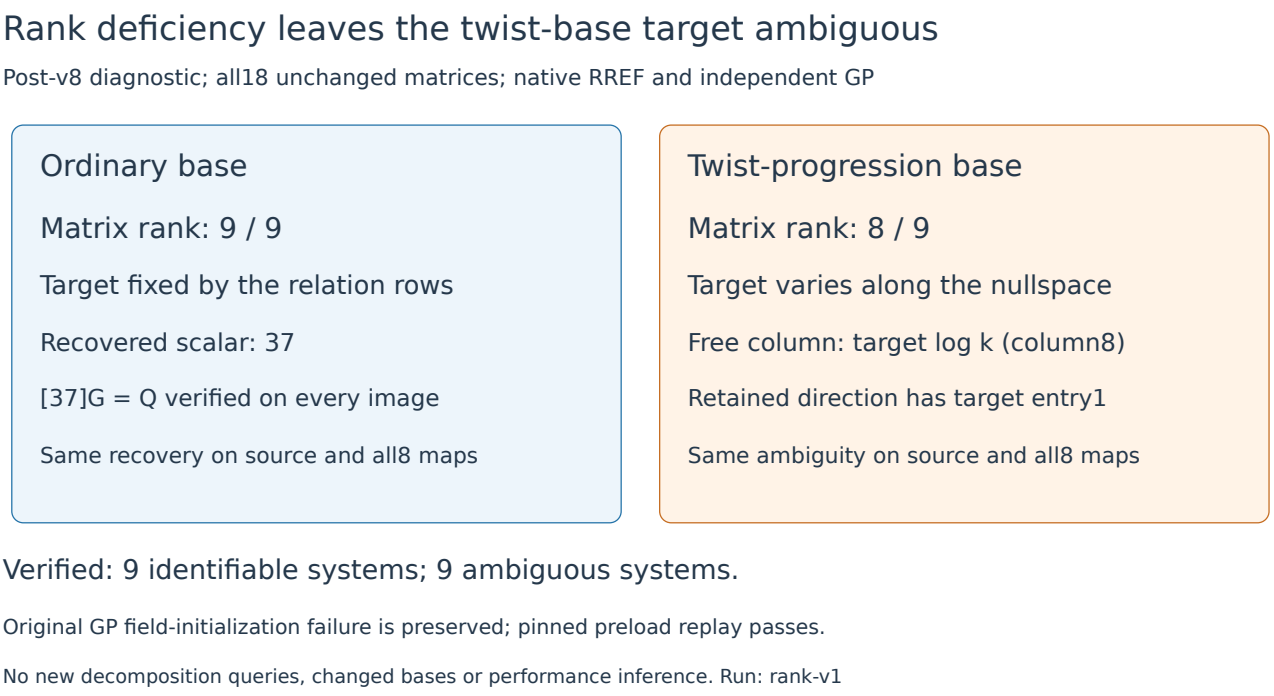

Figure 1: Target identifiability

Original source a3da15bda17804c75b48bbe487e48cb4ae5cb560; protocol fdf1d7d301ad748e ee6220c080f8bb15d118a361985496874075c757a85aaea1. The native run passes. The first GP process returns0 and prints its final marker but emits errors: the p assignment was parsed inside the checker definition, so the first matrix used a symbolic field. **That reference fails strict verification.** All bytes and the faithful failed-row export remain.

Before the second GP run, amendment/protocolv2 pins an explicit p823 preload. The main GP script, native output, matrix entries, ranks and expected flags are unchanged. The corrected replay has empty stderr, exact marker and exit0. Its protocol hash is 8c34026c04da1c26a9ff63735245a515 069723d61f255eeee098e4ff080bbf82. Both original and corrected exports pass their complete schema/protocol/file binding. They reference the **same single native attempt**, plus the respective actual GP attempt; do not pool them as four independent executions. Across both exports there are three actual processes: native1, GP2, with the first GP verification failed. A separate GP control initially used a vector in place of a one-row matrix; that failure and corrected passing control remain.

The original failed-row serializer used the wrong outcome class and its checker rejection is preserved;

the corrected export still says verification error/statusfailed, with outcomecompleted meaning the process finished. No failure is converted to a verified result. All clocks remain null/unadmitted.

Bounds:300s per native/GP process,4GiB,CPU0,one thread,L0 managed Xeon Platinum8573C/KVM. Primary PARI help and existing native elimination source were read; this is elementary linear algebra, not a novel attack. The post-v8 motivation is disclosed; it is not a pre-v8 prediction or a fresh target sample. No cost, timing, NIST or complete-IC improvement follows.

The raw archive retains all three actual processes and their receipts: 2,712 bytes,11 members, SHA256 80052a20f9d52aa5bad20dc1f1295bd878f76b24c290f7bcee78310a96fdde34. See reproduction, admission, failed original export, corrected replay export, amendment and archive verification.

This identifies which information is missing from this particular short progression base. The next distinct questions are whether another frozen target-blind base fixes the rank defect and whether any certified quotient reduces fully charged decomposition cost. Neither is measured here. Existing performance graphs stay unchanged; this diagnostic adds a correctness figure.